

ERIC ZHONG

☎ 267-481-6608 ✉ eric@erictzhong.com 🔗 linkedin.com/in/erictzhong 🐙 github.com/etzhong 🌐 erictzhong.com

EDUCATION

University of Southern California

GPA: 4.00

M.S. Electrical Engineering (Computer Architecture and Chip Design Specialization)

Graduated May 2025

B.S. Computer Engineering and Computer Science

(B.S. + M.S.)

Honors: Top of Class, Presidential Merit Scholarship, Provost Research Fellowship, Engineering Honors Fellowship

Relevant Coursework

CS: Algorithms and Computing Theory, Operating Systems, Compilers, Networking, C++, Data Structures and OOP, Software Engineering

ECE: ASIC Design and Tapeout, Computer Microarchitecture, Computer Systems Architecture, VLSI Circuit Design, System-on-Chip (SoCs), FPGAs, Digital System Design, Network Processor Design, Parallel Computing Systems, Embedded and Distributed Systems

Math: Probability and Statistics, Linear Algebra, Discrete Math, Multivariable Calculus, Differential Equations

EXPERIENCE

NVIDIA

Santa Clara, CA

GPU Architect

June 2025—Present

- Driving hardware design, architecture specification, **C++** simulation, and functional verification for NVIDIA's next-generation Feynman GPUs. Owned IPs include host **DMA engines** and work schedulers, **PCIe/CXL** endpoints, and interrupt controllers
- Writing custom GPU firmware in **C** and **NV-RISC-V assembly** for on-chip dynamic CUDA work scheduling and dependency resolution—enabling low-latency GPU kernel launch and improved performance with CUDA graph workflows
- Developing simulation tools in **C++** and **C** within NVIDIA's proprietary resman kernel driver and MODS testing framework for architectural studies and bring-up of new GPU hardware + software features
- Creating a new development infrastructure in **Python** and **Perl** for discovery of hardware device information by driver software
- Performing functional and performance testing using a **Python DSL** across **C++** simulation, **RTL**, **FPGA** emulation, and silicon
- Engineering a new **AI agent** for **DFD/DFT**—enabling autonomous debug and RCA of functional failures on production silicon

GPU Architecture Intern

May 2024—August 2024

- Architected and developed a **C++** functional model (fmodel) for a new GPU hardware IP enabling low-latency, real-time interrupt processing and work scheduling—facilitating RTL and kernel driver development for NVIDIA's flagship Rubin GPUs
- Built and deployed a new lightweight unit testing infrastructure using **C++** and **gtest** within NVIDIA's proprietary simulation environment—accelerating functional and performance testing by up to a 20x speedup over existing fullchip testing workflows
- Refactored legacy **C++/SystemC** fmodel with **C++20** source to improve core sim performance and future code maintainability

USC Photonic Computing Lab (Opticore)

Los Angeles, CA

Research Fellow

August 2022—May 2024

- Co-invented and published a novel teraFLOP optical tensor processor with a 25 fJ/Op energy efficiency—offering up to a 40x efficiency improvement over current state-of-the-art electronic processors for machine learning workflows ([Publication Link](#))
- Designed and verified a high-throughput (10+ GSPS), low-latency (< 100 ns) **FPGA**-based waveform generation and signal processing module using **SystemVerilog** and **C++**. Synthesized design onto an AMD Zynq UltraScale+ RFSoc using **Vivado**
- Developed a custom **Python** library for host interfacing with lab **FPGAs** along **100G Ethernet** to facilitate optical experiments

IQVIA

Philadelphia, PA

Software Engineering Intern

May 2022—August 2022

- Engineered a high-throughput data pipeline using **Scala** and **Python** to process incoming market and research data feeds for 300+ healthcare products—aggregating and consolidating 40+ independent data sources for uniform ingestion and analysis
- Prototyped several multivariate regression models in **Python** using **Jupyter Notebook** for market forecasting across 20+ sectors—achieving an aggregate performance of <5% SMAPE across a 3-month production testing window
- Built and deployed a **CI/CD** pipeline for **Scala microservices** backed by **Oracle DB**—automating pre-prod testing and SQA

PROJECTS

Cryptographic Accelerator ASIC

- Developed architecture and **SystemVerilog RTL** for a power and area-optimized ASIC for efficiently computing variable delay functions (VDFs)—achieving 30+ million 256-bit hash computations per second with a module size of 105 um x 105 um
- Taped out design using Intel's proprietary 16nm FinFET process in February 2025. Design brought up in November 2025

NetFPGA SmartNIC

- Designed a high-throughput network switch for smart packet routing and preprocessing by payload content on a **NetFPGA-1G**
- Architected core IP and wrote **Verilog RTL** for a multi-core NPU, custom network routing logic, and a software-programmable memory interface—achieving 2.5M packet inferences per second. Verified and synthesized design using **Verilator** and **Vivado**

Multithreaded ARM Processor

- Designed a 32-bit, 4-way multithreaded, 8-stage pipelined processor in **Verilog** running a custom ARMv4-compatible ISA
- Synthesized and implemented **RTL** onto an AMD Virtex **FPGA** using **Vivado**. Verified design using **Verilator** and **cocotb**
- Built a custom Clang-based **compiler** toolchain and backend for converting **C** code into hardware-executable **assembly**

SKILLS

Languages: C++, C, SystemVerilog, Verilog, Python, SystemC, Assembly, Rust, VHDL, Bash, Perl, Tcl

Technologies: FPGAs, Vivado, Vitis, Verilator, cocotb, ModelSim, Icarus Verilog, Quartus, CUDA, LLVM, Linux, Git, Perforce Helix

Domains: Hardware + Software Co-design, Computer Architecture, Systems Software, Open-Source Silicon, GPUs